

# Pattern Recognition

(EE5907R)

Jiashi FENG

Email: [elefjia@nus.edu.sg](mailto:elefjia@nus.edu.sg)

# Outlines

- Unsupervised Feature Extraction (PCA, NMF,...)
- Supervised Feature Extraction (LDA, GE, ...)
- Clustering and Applications
- **Gaussian Mixture Model and Boosting**
- Support Vector Machine
- Deep Learning

# Generative vs. Discriminative

- We want to classify the data  $x$  into labels  $y$ . A generative model learns the joint probability distribution  $p(x, y)$  and a discriminative model learns the conditional probability distribution  $p(y|x)$ .
- Suppose we have the following data in the form  $(x, y)$ : (1,0), (1,0), (2,0), (2, 1)

- $p(x, y)$  is

	$y = 0$	$y = 1$
$x = 1$	$\frac{1}{2}$	0
$x = 2$	$\frac{1}{4}$	$\frac{1}{4}$


- $p(y|x)$  is

	$y = 0$	$y = 1$
$x = 1$	1	0
$x = 2$	$\frac{1}{2}$	$\frac{1}{2}$

# Generative vs. Discriminative



- **Generative model**

(The artist)




- **Discriminative model**

(The lousy painter)




# Generative Models

- Gaussian Mixture Model and other mixture model




- Hidden Markov Model



# Discriminative Models

## Nearest neighbor




Shakhnarovich, Viola, Darrell 2003

Berg, Berg, Malik 2005

...

## Neural Networks




LeCun, Bottou, Bengio, Haffner 1998

Rowley, Baluja, Kanade 1998

...

## Support Vector Machines




Guyon, Vapnik

Heisele, Serre, Poggio, 2001

...

## Boosting



# Generative: Gaussian Mixture Model (GMM)

# Mixture Models

- Formally a Mixture Model is the weighted sum of a number of probability density functions (pdfs) where the weights are determined by a distribution,  $\pi$

$$p(x) = \pi_1 f_1(x) + \pi_2 f_2(x) + \dots + \pi_K f_K(x)$$

where  $\sum_{i=1}^K \pi_i = 1$

$$p(x) = \sum_{i=1}^K \pi_i f_i(x)$$



# Gaussian Mixture Models

- GMM: the weighted sum of a number of **Gaussians** where the weights are determined by a distribution,  $\pi$


$$p(x) = \pi_1 N(x|\mu_1, \Sigma_1) + \pi_2 N(x|\mu_2, \Sigma_2) + \dots + \pi_K N(x|\mu_K, \Sigma_K)$$

where  $\sum_{i=1}^K \pi_i = 1$

$$p(x) = \sum_{i=1}^K \pi_i N(x|\mu_i, \Sigma_i)$$


# Gaussian Mixture Models

- Rather than identifying clusters by “nearest” centroids
- Fit a Set of  $K$  Gaussians to the **unlabeled** data
- Maximum Likelihood over a mixture model




# The GMM Assumption

- There are  $K$  components. The  $i$ 'th component is called  $\omega_i$
- Component  $\omega_i$  has an associated mean vector  $\mu_i$



# The GMM Assumption

- There are  $K$  components. The  $i$ 'th component is called  $\omega_i$
- Component  $\omega_i$  has an associated mean vector  $\mu_i$
- Each component generates data from a Gaussian model with mean  $\mu_i$  and covariance matrix  $\sigma^2 I$




# The GMM Assumption

- There are  $K$  components. The  $i$ 'th component is called  $\omega_i$
- Component  $\omega_i$  has an associated mean vector  $\mu_i$
- Each component generates data from a Gaussian model with mean  $\mu_i$  and covariance matrix  $\sigma^2 I$

Assume that each data point is generated according to the following recipe:

1. Pick a component at random. Choose component  $i$  with probability  $P(\omega_i)$ .




# The GMM Assumption

- There are  $K$  components. The  $i$ 'th component is called  $\omega_i$
- Component  $\omega_i$  has an associated mean vector  $\mu_i$
- Each component generates data from a Gaussian model with mean  $\mu_i$  and covariance matrix  $\sigma^2 I$

Assume that each data point is generated according to the following recipe:

1. Pick a component at random. Choose component  $i$  with probability  $P(\omega_i)$ .
2. Data point  $\sim N(\mu_i, \sigma^2 I)$




# The General GMM Assumption

- There are  $K$  components. The  $i$ 'th component is called  $\omega_i$
- Component  $\omega_i$  has an associated mean vector  $\mu_i$
- Each component generates data from a Gaussian model with mean  $\mu_i$  and covariance matrix  $\Sigma_i$

Assume that each data point is generated according to the following recipe:

1. Pick a component at random. Choose component  $i$  with probability  $P(\omega_i)$ .
2. Data point  $\sim N(\mu_i, \Sigma_i)$




# The EM Algorithm

- **Expectation-maximization (EM)** is a method for finding **maximum likelihood** (or maximum a posteriori) estimate of parameter(s) in statistical model, where the model depends on **unobserved latent variables**.

Latent variables are the key properties for EM.



# The EM Algorithm

- **EM** is an **iterative** method which alternates between performing an **Expectation (E)** step and a **Maximization (M)** step
    - **E-step** computes the expectation of the log-likelihood evaluated using the current estimated distributions for the latent variables based on the parameters inferred from previous step
    - **M-step** computes parameters maximizing the expected log-likelihood from the E-step. These parameter-estimates are then used to determine the distribution of the latent variables in the next E-step.
- 

Latent variables  
become  
constants here.

# Simple Example

Let events be “grades in a class”

$$w_1 = \text{Gets an A} \quad P(A) = \frac{1}{2}$$

$$w_2 = \text{Gets a B} \quad P(B) = \mu$$

$$w_3 = \text{Gets a C} \quad P(C) = 2\mu$$

$$w_4 = \text{Gets a D} \quad P(D) = \frac{1}{2} - 3\mu$$

(Note  $0 \leq \mu \leq 1/6$ )

Assume we want to estimate  $\mu$  from data. In a given class, there were

a A's  
b B's  
c C's  
d D's

What's the maximum likelihood estimate of  $\mu$  given a, b, c, d?

# Trivial Statistics

$$P(A) = \frac{1}{2} \quad P(B) = \mu \quad P(C) = 2\mu \quad P(D) = \frac{1}{2} - 3\mu$$

$$P(a, b, c, d | \mu) = K \left(\frac{1}{2}\right)^a (\mu)^b (2\mu)^c \left(\frac{1}{2} - 3\mu\right)^d$$

$$\log P(a, b, c, d | \mu) = \log K + a \log \frac{1}{2} + b \log \mu + c \log 2\mu + d \log \left(\frac{1}{2} - 3\mu\right)$$

FOR MAX LIKE  $\mu$ , SET  $\frac{\partial \text{Log} P}{\partial \mu} = 0$

$$\frac{\partial \text{Log} P}{\partial \mu} = \frac{b}{\mu} + \frac{2c}{2\mu} - \frac{3d}{1/2 - 3\mu} = 0$$

$$K = \frac{(a + b + c + d)!}{a!b!c!d!}$$

Gives max like

$$\mu = \frac{b + c}{6(b + c + d)}$$

So if class got

A	B	C	D
14	6	9	10

$$\text{Max like } \mu = \frac{1}{10}$$

# Same Problem with Latent Information

Someone tells us that

Number of High grades (A's + B's) =  $h$

Number of C's =  $c$

Number of D's =  $d$

What is the max likelihood estimate of  $\mu$  now?

REMEMBER

$$P(A) = \frac{1}{2}$$

$$P(B) = \mu$$

$$P(C) = 2\mu$$

$$P(D) = \frac{1}{2} - 3\mu$$

$$\log P(h, c, d \mid \mu, b) = \log K(h-b, b, c, d) + (h-b) \log \frac{1}{2} + b \log \mu + c \log 2\mu + d \log (\frac{1}{2} - 3\mu)$$

latent variable

# Same Problem with Latent Information

Someone tells us that

Number of High grades (A's + B's) =  $h$

Number of C's =  $c$

Number of D's =  $d$

What is the max likelihood estimate of  $\mu$  now?

We can answer this question circularly:

REMEMBER

$$P(A) = \frac{1}{2}$$

$$P(B) = \mu$$

$$P(C) = 2\mu$$

$$P(D) = \frac{1}{2} - 3\mu$$

## EXPECTATION

If we know the value of  $\mu$  we could compute the expected value of  $b$

Since the ratio  $a:b$  should be the same as the ratio  $\frac{1}{2} : \mu$

$$E_{\mu}(b) = \frac{\mu}{\frac{1}{2} + \mu} h$$

## MAXIMIZATION

If we know the expected values of  $b$  we could compute the maximum likelihood value of  $\mu$

$$\mu = \frac{E_{\mu}(b) + c}{6(E_{\mu}(b) + c + d)}$$

Already computed as in slide #19

# EM for This Problem

REMEMBER

$$P(A) = \frac{1}{2}$$

$$P(B) = \mu$$

$$P(C) = 2\mu$$

$$P(D) = \frac{1}{2} - 3\mu$$

We begin with a guess for  $\mu$

We iterate between **EXPECTATION** and **MAXIMIZATION**  
to improve our estimates of  $b$  and  $\mu$ .

Define  $\mu(t)$  the estimate of  $\mu$  on the  $t$ 'th iteration

$b(t)$  the estimate of  $b$  on  $t$ 'th iteration


$\mu(0)$  = initial guess

$$b(t) = \frac{\mu(t)h}{\frac{1}{2} + \mu(t)} = E[b \mid \mu(t)]$$



$$\mu(t+1) = \frac{b(t) + c}{6(b(t) + c + d)}$$

= max like est of  $\mu$  given  $b(t)$



**Continue iterating until converged.**

**Good news: Converging to local optimum is assured.**


**Bad news: We have “local” optimum.**

# EM Convergence

- Convergence proof based on fact that  $\text{Prob}(\text{data} \mid \mu)$  must increase or remain same between each iteration [NOT OBVIOUS, BUT NOT STUDY HERE]
- But it can never exceed 1
- So it must therefore converge

In our example, suppose  
we had

$h = 20$   
 $c = 10$   
 $d = 10$   
 $\mu(0) = 0$



t	$\mu(t)$	b(t)
0	0	0
1	0.0833	2.857
2	0.0937	3.158
3	0.0947	3.185
4	0.0948	3.187
5	0.0948	3.187
6	0.0948	3.187

# Back to Learning of GMM

Remember:

We have unlabeled data  $\mathbf{x}_1 \mathbf{x}_2 \dots \mathbf{x}_N$

We know there are  $K$  components

We know  $P(\omega_1) P(\omega_2) P(\omega_3) \dots P(\omega_k), \sigma$

We don't know  $\boldsymbol{\mu}_1 \boldsymbol{\mu}_2 \dots \boldsymbol{\mu}_k$

Hidden variables  $z_k^n$ , indicating which component  $k$  the datum  $n$  is sampled from



# Compute Likelihood

- We define:

$$\pi_i = P(\omega_i) \quad \text{where} \quad \sum_i \pi_i = 1$$

$$z_i = p(\omega_i|x) = \frac{P(\omega_i)p(x|\omega_i)}{\sum_{j=1}^K P(\omega_j)p(x|\omega_j)}$$

$$z_k^n = p(\omega_k|x_n)$$

- Identify a likelihood function

$$\begin{aligned} p(x_1, \dots, x_N|\pi, \mu) &= \prod_{n=1}^N p(x_n|\pi, \mu) \\ &= \prod_{n=1}^N \sum_{k=1}^K p(x_n|\omega_k, \mu_k) P(\omega_k) \end{aligned}$$

$x_n$ 's were drawn independently

# Maximum Likelihood over a GMM

- Identify a log-likelihood function

$$\ln p(x_1, \dots, x_n | \pi, \mu) = \sum_{n=1}^N \ln \left[ \sum_{k=1}^K p(x_n | \omega_k, \mu_k) P(\omega_k) \right]$$

- Compute and set partials to 0

$$\frac{\partial \ln p(x_1, \dots, x_n | \pi, \mu)}{\partial \mu_k} = \sum_{n=1}^N \frac{1}{p(x_n | \pi, \mu)} \frac{\partial \sum_{k=1}^K N(x_n | \mu_k) P(\omega_k)}{\partial \mu_k}$$

$$p(x_n | \pi, \mu) = \sum_{k=1}^K p(x_n | \omega_k, \mu_k) P(\omega_k)$$

$$= \sum_{n=1}^N \frac{P(\omega_k)}{p(x_n | \pi, \mu)} \frac{\partial N(x_n | \mu_k)}{\partial \mu_k}$$

$$= \sum_{n=1}^N \frac{P(\omega_k) N(x_n | \mu_k)}{p(x_n | \pi, \mu)} \frac{\partial \ln N(x_n | \mu_k)}{\partial \mu_k}$$

$$\frac{\partial \ln f(x)}{\partial x} = \frac{1}{f(x)} \frac{\partial f(x)}{\partial x}$$

see next slide

# Maximum Likelihood over a GMM

$$\frac{\partial \ln p(x_1, \dots, x_n | \pi, \mu)}{\partial \mu_k} = \sum_{n=1}^N p(\omega_k | x_n) \frac{\partial \ln \exp\left(-\frac{1}{2\sigma^2} (x_n - \mu_k)^2\right)}{\partial \mu_k}$$

$$= \sum_{n=1}^N z_k^n \left( -\frac{1}{2\sigma^2} \frac{\partial (x_n - \mu_k)^2}{\partial \mu_k} \right)$$

$$= \sum_{n=1}^N z_k^n \frac{x_n - \mu_k}{\sigma^2} = 0$$

set partials to 0



$$\mu_k = \frac{\sum_{n=1}^N z_k^n x_n}{\sum_{n=1}^N z_k^n}$$

# EM for General GMMs

We don't know  $P(\omega_1), P(\omega_2), \dots, P(\omega_K), \mu_1, \mu_2, \dots, \mu_K, \Sigma_1, \Sigma_2, \dots, \Sigma_K$


Similarly, after compute the log likelihood and take partials to 0, we have

$$\mu_k = \frac{\sum_{n=1}^N z_k^n x_n}{\sum_{n=1}^N z_k^n}$$

$$\Sigma_k = \frac{\sum_{n=1}^N z_k^n (x_n - \mu_k)(x_n - \mu_k)^T}{\sum_{n=1}^N z_k^n}$$

$$\pi_k = \frac{\sum_{n=1}^N z_k^n}{N}$$

# Summary: EM for GMMs

- Initialize the parameters
    - Evaluate the log likelihood
  - Expectation-step: Compute the expectation
  - Maximization-step: Re-estimate Parameters
    - Evaluate the log likelihood
    - Check for convergence
- 

# EM for GMMs

- E-step: Compute “expected” classes of all data points for each class

$$z_k^n = \frac{\pi_k N(x_n | \mu_k, \Sigma_k)}{\sum_{j=1}^K \pi_j N(x_n | \mu_j, \Sigma_j)}$$

where  $\pi_k = p(\omega_k)$

# EM for GMMs

- M-Step: Re-estimate Parameters


$$\mu_k^{new} = \frac{\sum_{n=1}^N z_k^n x_n}{\sum_{n=1}^N z_k^n}$$

$$\Sigma_k^{new} = \frac{\sum_{n=1}^N z_k^n (x_n - \mu_k^{new})(x_n - \mu_k^{new})^T}{\sum_{n=1}^N z_k^n}$$

$$\pi_k^{new} = p(\omega_k)^{new} = \frac{\sum_{n=1}^N z_k^n}{N}$$


Latent variables  
become  
constants here.

# Gaussian Mixture Model Example: Start







# After 1st iteration




# After 2nd iteration




# After 3rd iteration




# After 4th iteration




# After 5th iteration



# After 6th iteration



# After 20th iteration




# Relationship to K-means


- K-means makes **hard** decisions.
  - Each data point gets assigned to a single cluster.
- GMM makes **soft** decisions.
  - Each data point yields a posterior
- Potential problem:
  - Incorrect number of Mixture Components



# Incorrect Number of Gaussians




# Incorrect Number of Gaussians




# GMM for Classification

- Train universal GMM, and then adapt it for individual class, and finally do classification
  - Widely used in speech recognition
- Note that we can initiate GMM by using K-means

# Another Application





**Discriminative: Boosting**


# Example Task

Object detection and recognition is formulated as a classification problem.


The image is partitioned into a set of overlapping windows

... and a decision is taken at each window about if it contains a target object or not.

Where are the screens?




Bag of image patches



In some feature space

# Formulation

- Formulation: binary classification



- Classification function

$\hat{y} = F(x)$  where  $F(x)$  belongs to some family of functions

- Minimize misclassification error

# Why Boosting?

- A simple algorithm for learning robust classifiers
  - Freund & Shapire, 1995
  - Friedman, Hastie, Tibshirani, 1998
- Provides efficient algorithm for sparse visual feature selection
  - *Tieu & Viola, 2000*
  - *Viola & Jones, 2003*
- Easy to implement, not requires external optimization tools



# Boosting

- Defines a classifier using an additive model:

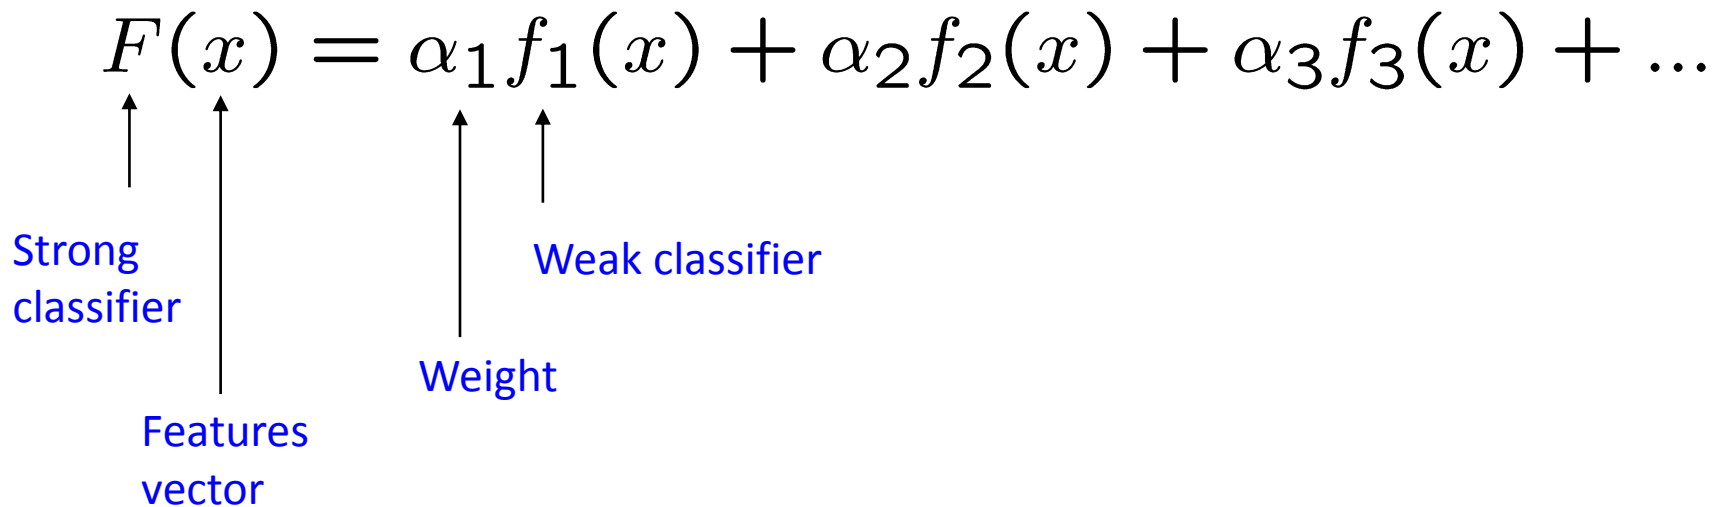
$$F(x) = \alpha_1 f_1(x) + \alpha_2 f_2(x) + \alpha_3 f_3(x) + \dots$$


Diagram illustrating the components of the additive model equation:

- $F(x)$ : Strong classifier
- $x$ : Features vector
- $\alpha_1$ : Weight
- $f_1(x)$ : Weak classifier

# Boosting

- Defines a classifier using an additive model:

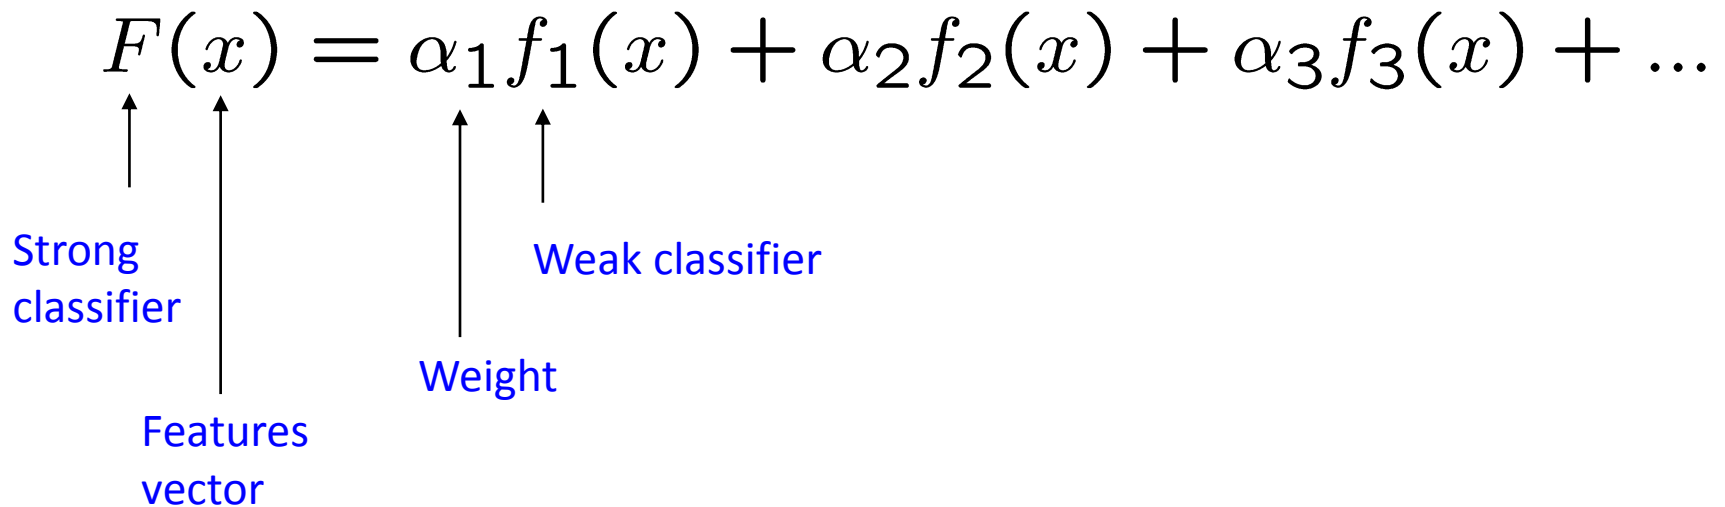
$$F(x) = \alpha_1 f_1(x) + \alpha_2 f_2(x) + \alpha_3 f_3(x) + \dots$$


Diagram illustrating the components of the additive model equation:


- $F(x)$  is labeled as the **Strong classifier**.
- $x$  is labeled as the **Features vector**.
- $\alpha_1$  is labeled as the **Weight**.
- $f_1(x)$  is labeled as the **Weak classifier**.

- We need to define a family of weak classifiers

$f_k(x)$  from a family of weak classifiers

# Toy Example

- It is a sequential procedure:



Each data point has  
a class label:


$$y_t = \begin{cases} +1 & (\text{red circle}) \\ -1 & (\text{blue circle}) \end{cases}$$

and a weight:

$$w_t = 1$$

# Toy Example

Weak learners from the family of lines



Each data point has  
a class label:


$$y_t = \begin{cases} +1 & (\text{red circle}) \\ -1 & (\text{blue circle}) \end{cases}$$

and a weight:

$$w_t = 1$$

Weak classifier  $h \Rightarrow p(\text{error}) = 0.5$  it is at chance

# Toy Example



Each data point has  
a class label:

$$y_t = \begin{cases} +1 & (\text{red circle}) \\ -1 & (\text{blue circle}) \end{cases}$$


and a weight:

$$w_t = 1$$

This one seems to be the best

This is a '**weak classifier**': It performs slightly better than chance.

# Toy Example



Each data point has  
a class label:

$$y_t = \begin{cases} +1 & (\text{red circle}) \\ -1 & (\text{blue circle}) \end{cases}$$


**We update the weights:**

$$w_t \leftarrow w_t \exp\{-y_t F(x_t)\}$$

Current weak  
classifier

We set a new problem for which the current classifier performs at chance again

# Toy Example



Each data point has  
a class label:


$$y_t = \begin{cases} +1 & (\text{red circle}) \\ -1 & (\text{blue circle}) \end{cases}$$

**We update the weights:**

$$w_t \leftarrow w_t \exp\{-y_t F(x_t)\}$$

Similarly, we learn another weak classifier

# Toy Example



Reweighting

Each data point has  
a class label:


$$y_t = \begin{cases} +1 & (\text{red circle}) \\ -1 & (\text{blue circle}) \end{cases}$$

**We update the weights:**

$$w_t \leftarrow w_t \exp\{-y_t F(x_t)\}$$



# Toy Example



Each data point has  
a class label:


$$y_t = \begin{cases} +1 & (\text{red circle}) \\ -1 & (\text{blue circle}) \end{cases}$$

**We update the weights:**

$$w_t \leftarrow w_t \exp\{-y_t F(x_t)\}$$

Similarly, we learn another weak classifier

# Toy Example



The strong (non- linear) classifier is built as the combination of all the weak (linear) classifiers.

# Boosting

- For different cost function and minimization algorithm, the result is a different flavor of Boosting
- We shall introduce gentleBoosting
  - It is simple to implement and numerically stable.

# Boosting

Boosting fits the additive model

$$F(x) = f_1(x) + f_2(x) + f_3(x) + \dots$$


by minimizing the exponential loss

$$J(F) = \sum_{t=1}^N e^{-y_t F(x_t)}$$

↑            ↑  
Training samples

The exponential loss is a differentiable upper bound to the misclassification error.

# Exponential Loss



Squared error

$$J = \sum_{t=1}^N [y_t - F(x_t)]^2$$

Exponential loss

$$J = \sum_{t=1}^N e^{-y_t F(x_t)}$$

# Boosting

Sequential procedure. At each step  $m$  we add

$$F(x) \leftarrow F(x) + f_m(x)$$

to minimize the residual loss

$$(\phi_m) = \arg \min_{\phi} \sum_{t=1}^N J(y_t, F(x_t) + f(x_t; \phi))$$

Parameters of the  
weak classifier

Desired output

input



# gentleBoosting

- At each iteration:

We chose  $f_m(x)$  that minimizes the cost:

$$J(F + f_m) = \sum_{t=1}^N e^{-y_t(F(x_t) + f_m(x_t))}$$

Instead of doing exact optimization, gentle Boosting minimizes the **approximation** of the error:

$$J(F) \propto \sum_{t=1}^N \boxed{e^{-y_t F(x_t)}} (y_t - f_m(x_t))^2 \rightarrow$$

At each iterations we just need to solve a weighted least squares problem

↑  
Weights at this iteration




# Weak Classifiers

- The input is a set of weighted training samples  $(x, y, w)$
- Regression stumps: simple but commonly used in object detection.

$$f_m(x) = a[x_k < \theta] + b[x_k \geq \theta]$$

Four parameters:  $\phi = [a, b, \theta, k]$









# Features -> Weak Detectors

## Haar filters and integral image

Viola and Jones, ICCV 2001




The average intensity in the block is computed with four sums independently of the block size.



$$\text{Sum} = D - B - C + A$$

# Features -> Weak Detectors

For screen detection, we may collect a set of part templates from a set of training objects to build feature set.



# Example: Screen Detection


Feature  
output



# Example: Screen Detection


Feature  
output

Thresholded  
output




Weak 'detector'  
Produces many false alarms.

# Example: Screen Detection




# Example: Screen Detection



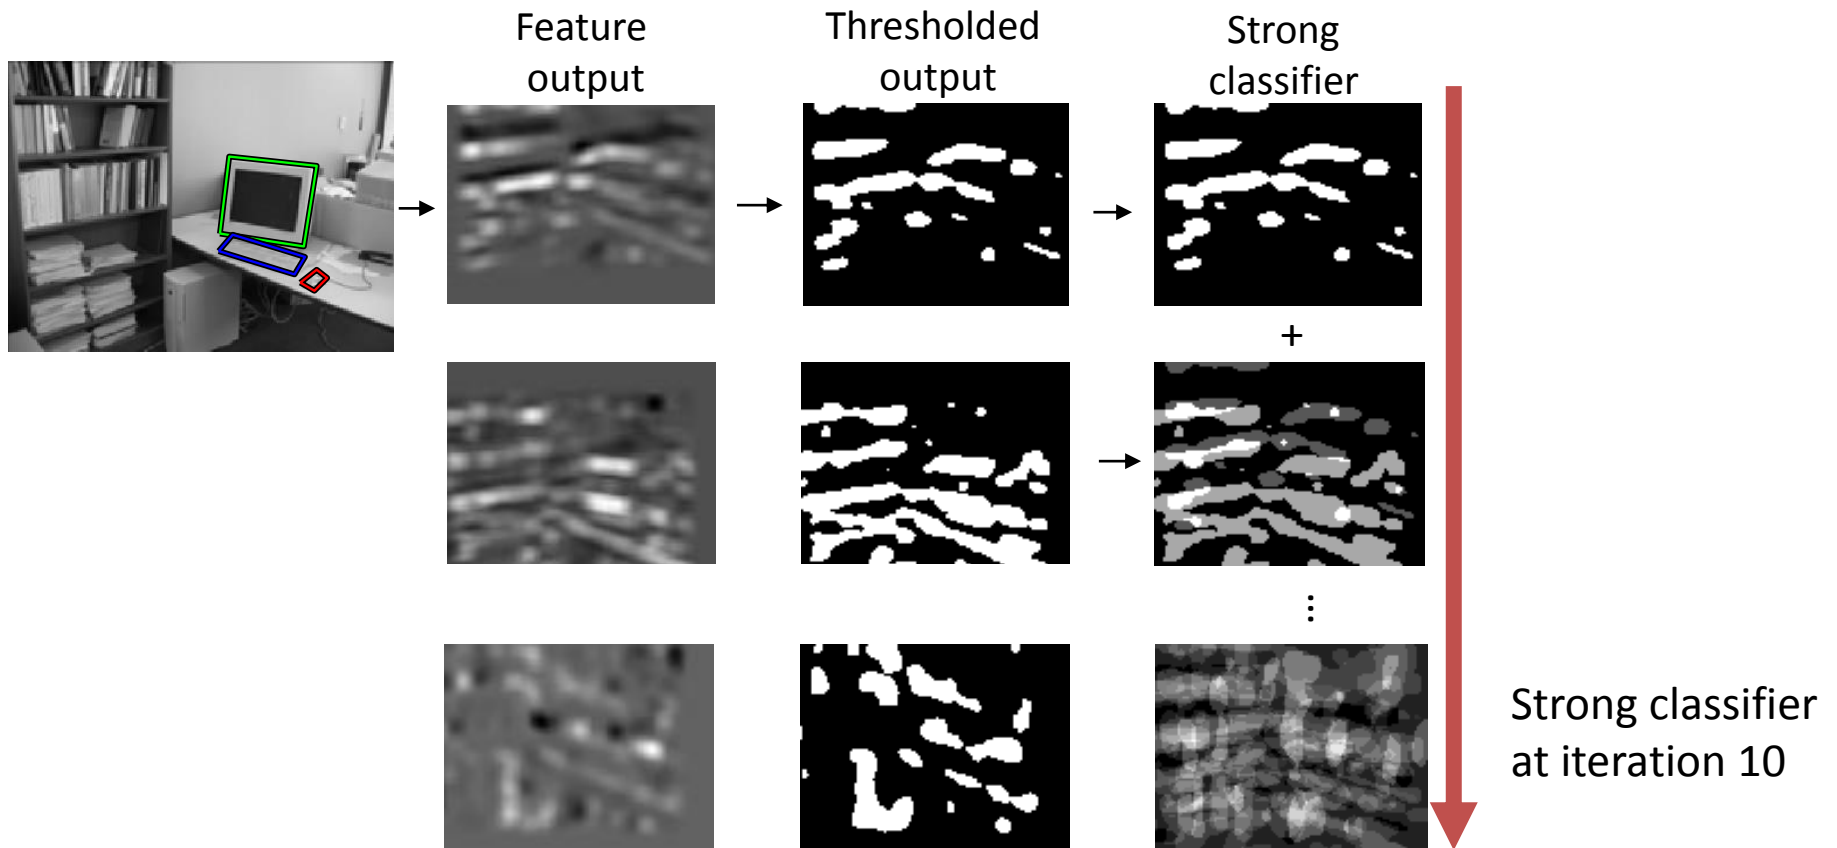
Second weak 'detector'

Produces a different set of false alarms.

# Example: Screen Detection

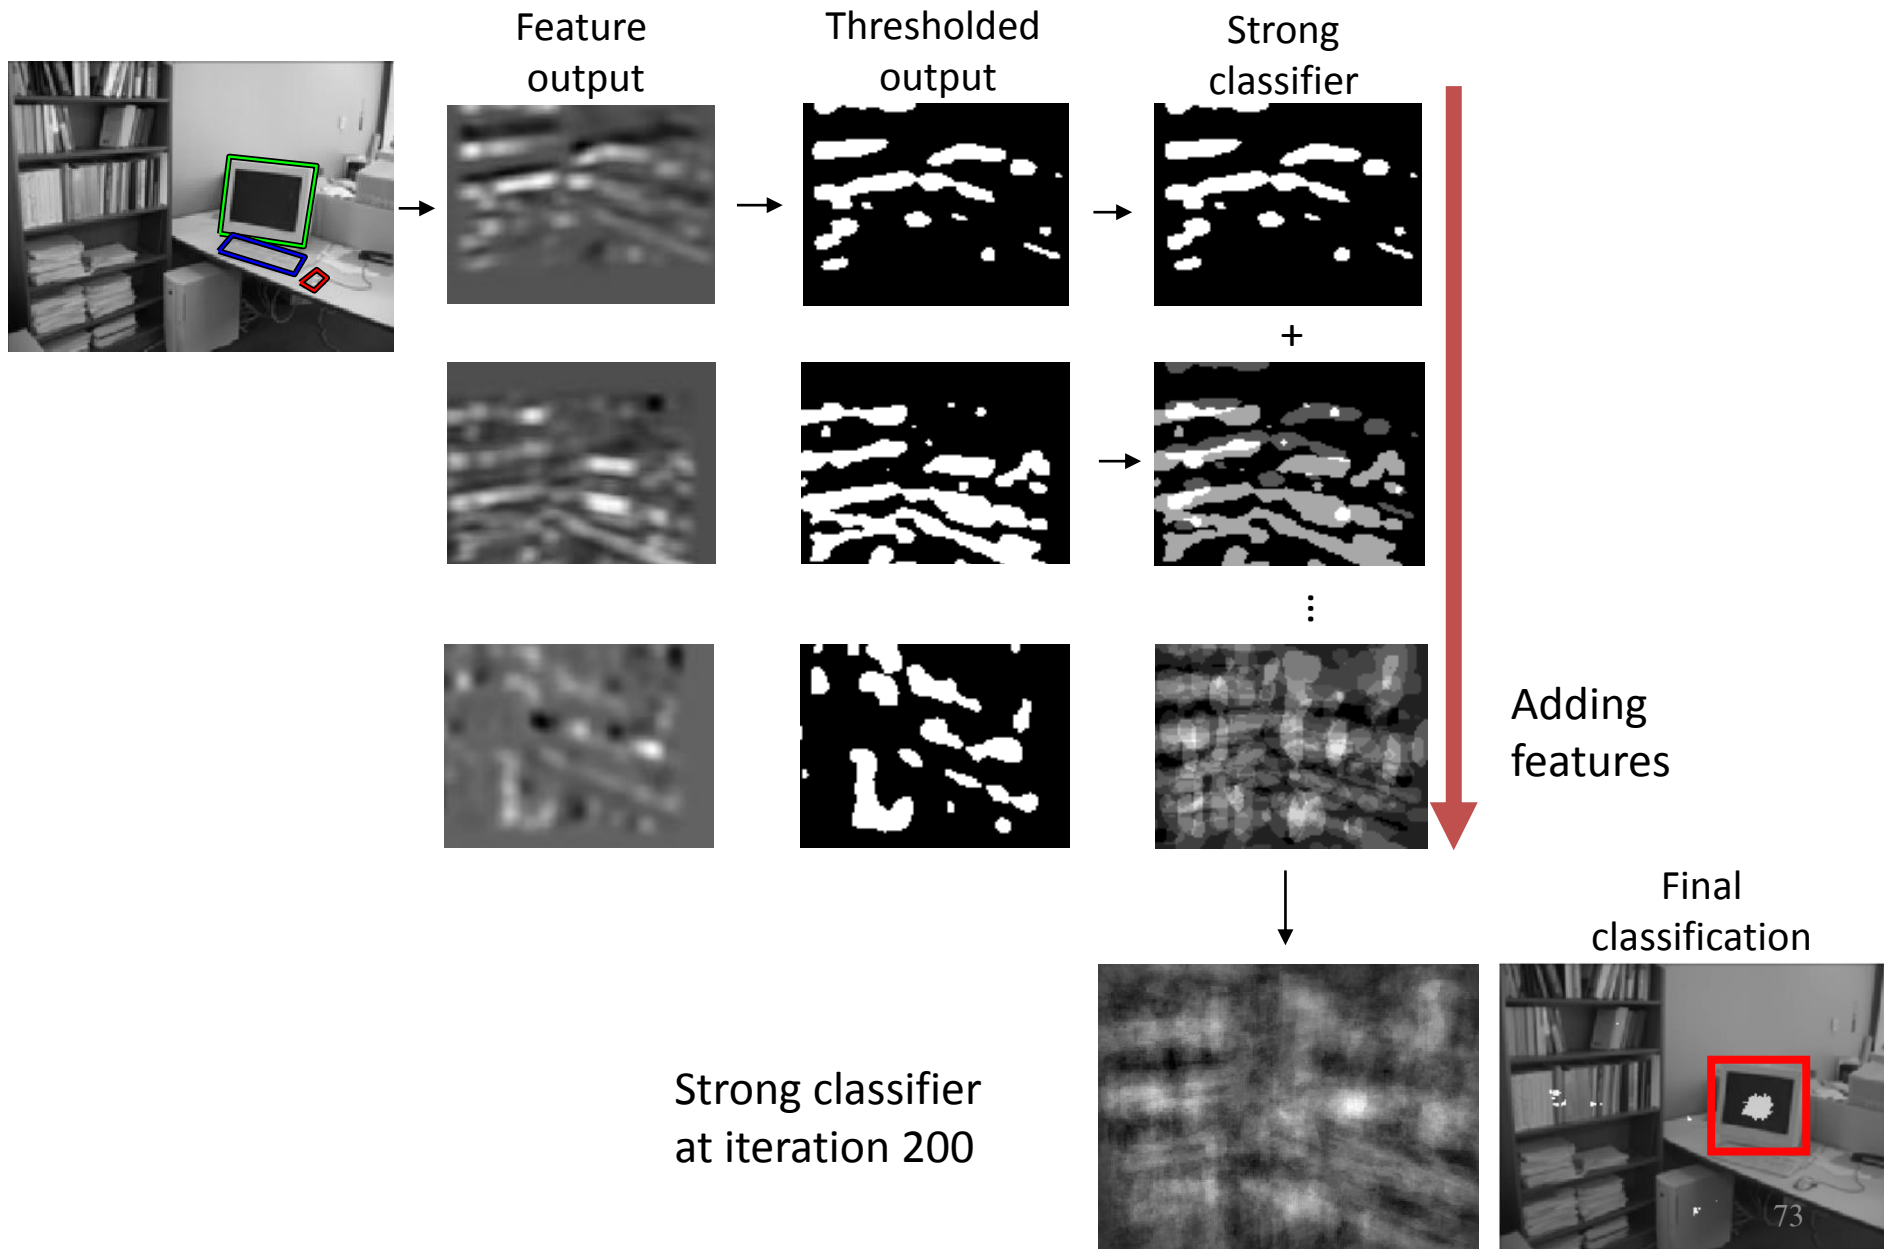


# Example: Screen Detection






# Example: Screen Detection




# Cascade of classifiers

What is the motivation: some negative samples may be rejected based on few features!




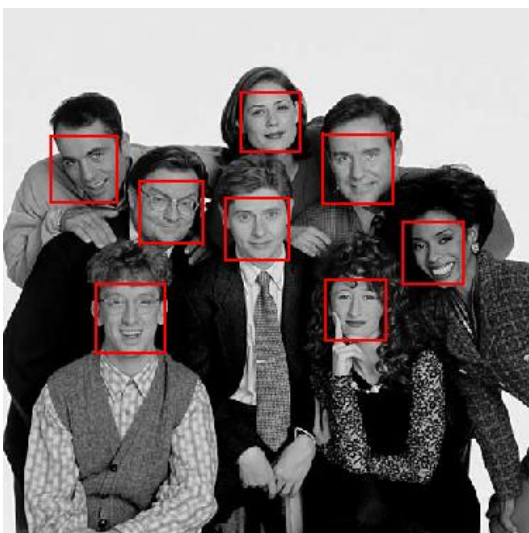


We want the complexity of the 3 features classifier with the performance of the 100 features classifier:



Select a threshold with high recall for each stage.

We increase precision using the cascade


# Output of Face Detector on Test Images



# Other detection tasks




Facial Feature Localization




Profile Detection

Male vs.  
female





# Profile Detection



# “Head in the coffee beans problem”

Can you find the head in this image?





# Weakness of Boosting

- Features are extracted at fixed positions, and thus not deformable (not perfect for deformable objects)
- No mechanism for handling occlusion
- Extension to “deformable model” + “and/or model”?



# Papers to Read and Study

- Friedman, Hastie, Tibshirani.

**Additive Logistic Regression: a Statistical View of Boosting** (1998). [Pdf](#)

- Robert E. Schapire.

**The boosting approach to machine learning: An overview.**

In D. D. Denison, M. H. Hansen, C. Holmes, B. Mallick, B. Yu, editors, *Nonlinear Estimation and Classification*. Springer, 2003.

[Postscript](#) or [gzipped postscript](#).

- Ron Meir and Gunnar Rätsch.

**An introduction to boosting and leveraging.**

In *Advanced Lectures on Machine Learning (LNAI2600)*, 2003. [Pdf](#).